Gaussian Process

Marshal Sinaga - marshal.sinaga@aalto.fi

Last update: March 25, 2024

This note aims to cover some materials on the Gaussian process. The primary references are Gaussian Process for Machine Learning by C. E. Rasmussen and CS-E4895 by Arno Solin.

1 Multivariate Normal Distribution

1.1 Linear transformation theorem for the multivariate normal distribution

Let x follow a multivariate normal distribution:

$$x \sim \mathcal{N}(\mu, \Sigma) \tag{1}$$

Then, any affine transformation of x is also multivariate normally distributed:

$$y = Ax + b \sim \mathcal{N}(A\mu + b, A\Sigma A^\top) \tag{2}$$

Proof:

The moment-generating function of random vector x is

$$M_x(t) = \mathbb{E}[\exp(t^T x)] \tag{3}$$

and therefore, the moment-generating function of the random vector y is given by

$$\begin{aligned} M_y(t) &= \mathbb{E}\left[\exp(t^T(Ax+b))\right] \\ &= \mathbb{E}[\exp(t^\top Ax)\exp(t^\top b)] \\ &= \exp(t^\top b)\mathbb{E}[\exp(t^\top Ax)] \\ &= \exp(t^\top b)M_x(A^\top t) \end{aligned} \tag{4}$$

The moment-generating function of the multivariate normal distribution is

$$M_x(t) = \exp(t^\top \mu + \frac{1}{2}t^\top \Sigma t) \tag{5}$$

and therefore, the moment-generating function of random vector y becomes

$$M_y(t) = \exp(t^\top(A\mu + b) + \frac{1}{2}t^\top A\Sigma A^\top t) \tag{6}$$

Since the moment-generating function and the probability density function of a random variable are equivalent, this demonstrates that y follows a multivariate normal distribution with mean $A\mu + b$ and covariance $A\Sigma A^\top$.

1.2 Marginal distribution of the multivariate normal distribution

Let x follow a multivariate normal distribution:

$$x \sim \mathcal{N}(\mu, \Sigma) \tag{7}$$

Then, the marginal distribution of any subset vector x_s is also a multivariate normal distribution.

$$x_s \sim \mathcal{N}(\mu_s, \Sigma_s) \tag{8}$$

where μ_s drops the irrelevant variables (the ones not in the subset, i.e., marginalized out) from the mean vector μ and Σ_s drops the corresponding rows and columns from the covariance matrix Σ.

Proof: Define an $m \times n$ subset matrix S such that $s_{ij} = 1$, if the $j-$th element in x_s corresponds to the $i-$th element in x, and $s_{ij} = 0$ otherwise. Then,

$$x_s = Sx \tag{9}$$

and we can apply the linear transformation theorem to give

$$x_s \sim \mathcal{N}(S\mu, S\Sigma S^\top) \tag{10}$$

Finally, we see that $S\mu = \mu_s$ and $S\Sigma S^\top = \Sigma_s$

1.3 Conditional distribution of the multivariate normal distribution

Let x follow a multivariate normal distribution

$$x \sim \mathcal{N}(\mu, \Sigma) \tag{11}$$

Then, the conditional distribution of any subset vector x_1, given the complement vector x_2, is also a multivariate normal distribution

$$\begin{aligned} \mu_{1|2} &= \mu_1 + \Sigma_{12}\Sigma_{22}^{-1}(x_2 - \mu_2) \\ \Sigma_{1|2} &= \Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21} \end{aligned} \tag{12}$$

with block-wise mean and covariance defined as:

$$\begin{aligned} \mu &= \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix} \\ \Sigma &= \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix} \end{aligned} \tag{13}$$

Proof: Without loss of generality, we assume that in parallel to 13,

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \tag{14}$$

where $x_1 \in \mathbb{R}^{n_1 \times 1}$, $x_2 \in \mathbb{R}^{n_2 \times 1}$, and $x \in \mathbb{R}^{n \times 1}$. The joint distribution of x_1 and x_2 is

$$x \sim \mathcal{N}(\mu, \Sigma) \tag{15}$$

Moreover, the marginal distribution of x_2 follows from 11 and 13 as

$$x_2 \sim \mathcal{N}(\mu_2, \Sigma_{22}) \tag{16}$$

According to conditional probability, it holds that

$$\begin{aligned} p(x_1|x_2) &= \frac{p(x_1, x_2)}{p(x_2)} \\ &= \frac{\mathcal{N}(\mu, \Sigma)}{\mathcal{N}(\mu_2, \Sigma_{22})} \end{aligned} \tag{17}$$

Using the probability density of multivariate-normal, this becomes

$$\begin{aligned} p(x_1|x_2) &= \frac{1/\sqrt{(2\pi)^n|\Sigma|}\exp\left(-\frac{1}{2}(x-\mu)^\top\Sigma^{-1}(x-\mu)\right)}{1/\sqrt{(2\pi)^{n_2}|\Sigma_{22}|}\exp\left(-\frac{1}{2}(x-\mu_2)^\top\Sigma_{22}^{-1}(x-\mu_2)\right)} \\ &= 1/\sqrt{(2\pi)^{n-n_2}}\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\exp\left(-\frac{1}{2}(x-\mu)^\top\Sigma^{-1}(x-\mu) + \frac{1}{2}(x-\mu_2)^\top\Sigma_{22}^{-1}(x-\mu_2)\right) \end{aligned} \tag{18}$$

Writing the inverse Σ as

$$\Sigma^{-1} = \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix} \tag{19}$$

and applying 13 to 18, we obtain:

$$\begin{aligned} p(x_1|x_2) = 1/\sqrt{(2\pi)^{n-n_2}}\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\exp(-\frac{1}{2}\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}\right)^\top \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix} \left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}\right) \\ + \frac{1}{2}(x-\mu_2)^\top\Sigma_{22}^{-1}(x-\mu_2)) \end{aligned} \tag{20}$$

Multiplying within 20, we have

$$\begin{aligned} p(x_1|x_2) = 1/\sqrt{(2\pi)^{n-n_2}}\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\exp(-\frac{1}{2}((x_1-\mu_1)^\top\Sigma^{11}(x_1-\mu_1) + 2(x_1-\mu_1)\Sigma^{12}(x_2-\mu_2) \\ + (x_2-\mu_2)^\top\Sigma^{22}(x_2-\mu_2)) + \frac{1}{2}(x-\mu_2)^\top\Sigma_{22}^{-1}(x-\mu_2)) \end{aligned} \tag{21}$$

where we have used the fact that $\Sigma^{12} = \Sigma^{21^\top}$, because Σ^{-1} is symmetric. The inverse of a block matrix is

$$\begin{bmatrix} A & B \\ C & D \end{bmatrix}^{-1} = \begin{bmatrix} (A-BD^{-1}C)^{-1} & -(A-BD^{-1}C)^{-1}BD^{-1} \\ -D^{-1}C(A-BD^{-1}C)^{-1} & D^{-1}+D^{-1}C(A-BD^{-1}C)^{-1}BD^{-1} \end{bmatrix} \tag{22}$$

Thus, the inverse of Σ^{-1} in 19 is

$$\begin{bmatrix}\Sigma_{11} & \Sigma_{12}\\ \Sigma_{21} & \Sigma_{22}\end{bmatrix}^{-1} = \begin{bmatrix}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1} & -(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\Sigma_{12}\Sigma_{22}^{-1}\\ -\Sigma_{22}^{-1}\Sigma_{21}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1} & \Sigma_{22}^{-1}+\Sigma_{22}^{-1}\Sigma_{21}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\Sigma_{12}\Sigma_{22}^{-1}\end{bmatrix} \tag{23}$$

Plugging this into 20, we have

$$\begin{aligned}p(x_1|x_2) =& \frac{1}{\sqrt{(2\pi)^{n-n_2}}}\cdot\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\cdot\\ & \exp\Bigg[-\frac{1}{2}\Big((x_1-\mu_1)^{\mathrm{T}}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}(x_1-\mu_1) - \\ & \qquad 2(x_1-\mu_1)^{\mathrm{T}}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)+\\ & \qquad (x_2-\mu_2)^{\mathrm{T}}\left[\Sigma_{22}^{-1}+\Sigma_{22}^{-1}\Sigma_{21}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\Sigma_{12}\Sigma_{22}^{-1}\right](x_2-\mu_2)\Big)\\ & +\frac{1}{2}\left((x_2-\mu_2)^{\mathrm{T}}\Sigma_{22}^{-1}(x_2-\mu_2)\right)\Bigg]\ .\end{aligned} \tag{24}$$

Eliminating some terms, we have

$$\begin{aligned}p(x_1|x_2) =& \frac{1}{\sqrt{(2\pi)^{n-n_2}}}\cdot\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\cdot\\ & \exp\Bigg[-\frac{1}{2}\Big((x_1-\mu_1)^{\mathrm{T}}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}(x_1-\mu_1) - \\ & \qquad 2(x_1-\mu_1)^{\mathrm{T}}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)+\\ & \qquad (x_2-\mu_2)^{\mathrm{T}}\Sigma_{22}^{-1}\Sigma_{21}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)\Big)\Bigg]\ .\end{aligned} \tag{25}$$

Rearranging the terms, we have

$$\begin{aligned}p(x_1|x_2) =& \frac{1}{\sqrt{(2\pi)^{n-n_2}}}\cdot\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\cdot\exp\Bigg[-\frac{1}{2}\cdot\\ & \left[(x_1-\mu_1)-\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)\right]^{\mathrm{T}}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\left[(x_1-\mu_1)-\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)\right]\Bigg]\\ =& \frac{1}{\sqrt{(2\pi)^{n-n_2}}}\cdot\sqrt{\frac{|\Sigma_{22}|}{|\Sigma|}}\cdot\exp\Bigg[-\frac{1}{2}\cdot\\ & \left[x_1-\left(\mu_1+\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)\right)\right]^{\mathrm{T}}(\Sigma_{11}-\Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21})^{-1}\left[x_1-\left(\mu_1+\Sigma_{12}\Sigma_{22}^{-1}(x_2-\mu_2)\right)\right]\Bigg]\end{aligned} \tag{26}$$

where we used the fact that $\Sigma_{21}=\Sigma_{12}^{\top}$. The determinant of a block matrix is

$$\begin{vmatrix}A & B\\ C & D\end{vmatrix} = |D|\cdot|A-BD^{-1}C|\ , \tag{27}$$

such that we have for Σ that

$$\begin{vmatrix}\Sigma_{11} & \Sigma_{12}\\ \Sigma_{21} & \Sigma_{22}\end{vmatrix} = |\Sigma_{22}| \cdot |\Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21}| \tag{28}$$

with this and $n - n_2 = n_1$, we finally arrive at

$$\begin{aligned} p(x_1|x_2) =& \frac{1}{\sqrt{(2\pi)^{n_1}|\Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21}|}} \cdot \exp\left[-\frac{1}{2}\cdot\right. \\ & \left.\left[x_1 - \left(\mu_1 + \Sigma_{12}\Sigma_{22}^{-1}(x_2 - \mu_2)\right)\right]^{\mathrm{T}} \left(\Sigma_{11} - \Sigma_{12}\Sigma_{22}^{-1}\Sigma_{21}\right)^{-1} \left[x_1 - \left(\mu_1 + \Sigma_{12}\Sigma_{22}^{-1}(x_2 - \mu_2)\right)\right]\right] \end{aligned} \tag{29}$$

which is the pdf of a multivariate normal distribution

$$p(x_1|x_2) = \mathcal{N}(x_1; \mu_{1|2}, \Sigma_{1|2}) \tag{30}$$

with mean $\mu_{1|2}$ and covariance $\Sigma_{1|2}$ given by 12.

2 The Marginal Likelihood

- Occam's razor: "When you have two competing models that produce similar predictions, the simpler, the better." The same concept goes for GP.
- The marginal likelihood $p(\mathbf{y}|\boldsymbol{\theta})$ implements a version of Occam's razor.
- Marginal likelihood for Gaussian likelihood

$$\begin{aligned} p(\mathbf{y}|\boldsymbol{\theta}) &= \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f}|\boldsymbol{\theta})d\mathbf{f} \\ &= \int \mathcal{N}(\mathbf{y}|\mathbf{f}, \sigma^2\mathbf{I})\mathcal{N}(\mathbf{f}|0, \mathbf{K})d\mathbf{f} \\ &= \mathcal{N}(\mathbf{y}|0, \sigma^2\mathbf{I} + \mathbf{K}) \end{aligned}$$

- Then

$$\log p(\mathbf{y}|\boldsymbol{\theta}) = \underbrace{-\frac{N}{2}\log(2\pi)}_{\text{constant}} \underbrace{-\frac{1}{2}\log|\sigma^2\mathbf{I} + \mathbf{K}|}_{\text{complexity penalty}} - \underbrace{\frac{1}{2}\mathbf{y}^\top(\sigma^2\mathbf{I} + \mathbf{K})^{-1}\mathbf{y}}_{\text{data fit}}$$

2.1 The Marginal Likelihood Computation

- In practice, we should avoid computing determinants and inverses.
- Step 1: Compute Cholesky factorization of $\mathbf{C} = \sigma^2\mathbf{I} + K$ such that $C = \mathbf{L}\mathbf{L}^\top$
- Step 2: Compute the log determinant as follows:

$$\log|\mathbf{C}| = \log|\mathbf{L}\mathbf{L}^\top| = \log|\mathbf{L}||\mathbf{L}^\top| = \log|\mathbf{L}|^2 = 2\log|\mathbf{L}| = 2\sum_{n=1}^{N}\log\mathbf{L}_{nn}$$

– Step 3: Compute quadratic term as follows

$$\mathbf{y}^\top \mathbf{C}^{-1}\mathbf{y} = \mathbf{y}^\top (\mathbf{L}\mathbf{L}^\top)^{-1}\mathbf{y} = \mathbf{y}^\top \mathbf{L}^{-\top}\mathbf{L}^{-1}\mathbf{y} = (\mathbf{L}^{-1}\mathbf{y})^\top \underbrace{(\mathbf{L}^{-1}\mathbf{y})}_{=\mathbf{v}} = \mathbf{v}^\top \mathbf{v}$$

– Step 4: Sum up components

$$\log p(\mathbf{y}|\boldsymbol{\theta}) = -\frac{N}{2}\log(2\pi) - \frac{1}{2}2\sum_{n=1}^{N}\log \mathbf{L}_{nn} - \frac{1}{2}\mathbf{v}^\top \mathbf{v}$$

- Note that we never compute the determinant or the inverse of $\mathbf{C}$ directly.

3 Kernel Theory

3.1 Hilbert Space

- A vector space $\mathcal{V}$ is a set of closed vectors under addition and scalar multiplication.
- If $\mathcal{V}$ is equipped with a norm $\|.\|_{\mathcal{V}} \in \mathbb{R}$, it is a norm space.
- A Hilbert space $\mathcal{H}$ is a complete inner product space, with inner product $\langle . \rangle_{\mathcal{H}}$ and induced norm $\|x\| = \sqrt{\langle x, x\rangle_{\mathcal{H}}}$.

3.2 Kernel Function and Reproducing Kernel Hilbert Space (RKHS)

- A function $k : \mathcal{X} \times \mathcal{X} \to \mathbb{R}$ is a kernel function if and only if there exists a Hilbert space $\mathcal{H}$ and a map $\phi : \mathcal{X} \to \mathcal{H}$ such that:

$$k(x, y) = \langle \phi(x), \phi(y) \rangle \tag{31}$$

for all $x, y \in \mathcal{X}$.

- Let $\phi : \mathcal{X} \to \mathbb{R}^{\mathcal{X}}$ and let us define:

$$k_x \coloneqq \phi(x) = k(x, .) \tag{32}$$

Therefore, we have $k_x(y) = k(x, y)$.

- Let $\mathcal{G}$ denote a vector space with span based on the images $\{k_x | x \in \mathcal{X}\}$, i.e.,

$$\{\mathcal{G} \coloneqq \sum_{i=1}^{m} \alpha_i k_{x_i} | \alpha_i \in \mathbb{R}, m \in \mathbb{N}, x_i \in \mathcal{X}\} \tag{33}$$

- By the definition of the kernel function, the inner product on $\mathcal{G}$ is defined as follows:

$$\langle k_x, k_y \rangle \coloneqq k(x, y) \tag{34}$$

Recall that $k_x = k(x, .)$, hence, $\langle k_x, k_y \rangle = \langle k(x, .), k(y, .) \rangle$.

- Therefore, for any $f, g \in \mathcal{G}$, with $f = \sum_i \alpha_i k_{x_i}$ and $g = \sum_j \beta_j k_{y_j}$, we have:

$$\langle f, g \rangle = \langle \sum_i \alpha_i k_{x_i}, \sum_j \beta_j k_{y_j} \rangle \tag{35}$$

$$= \sum_{ij} \alpha_i \beta_j \langle k_{x_i}, k_{y_j} \rangle \tag{36}$$

$$= \sum_{ij} \alpha_i \beta_j k(x_i, y_j) \tag{37}$$

- To make $\mathcal{G}$ a Hilbert space, we need to make it complete, i.e., ensure all Cauchy sequences converge.

Definition 1. *Let $\mathcal{H}$ be a Hilbert space of real function f defined on an index set $\mathcal{X}$. Then $\mathcal{H}$ is called a reproducing kernel Hilbert space endowed with an inner product $\langle ., . \rangle_{\mathcal{H}}$ if there exists a kernel function $k : \mathcal{X} \times \mathcal{X} \to \mathbb{R}$ with the following properties:*

1. *For every $x \in \mathcal{X}, k_x(y) = k(x, y)$ as function of $y \in \mathcal{X}$ belongs to $\mathcal{H}$, and*
2. *k has the reproducing property.*

- Reproducing property:

$$\langle k_x, f \rangle = \langle k_x, \sum_i \alpha_i k_{x_i} \rangle \tag{38}$$

$$= \sum_i \alpha_i \langle k_x, k_{x_i} \rangle = \sum_i k(x, x_i) = f(x) \tag{39}$$

- Moore-Aronszajn theorem: Given a kernel, there is a unique RKHS, Given an RKHS, there is a unique kernel.

3.3 Representer Theorem

Settings:

- We are given kernel k and denote the corresponding RKHS at $\mathcal{H}$.
- We want to learn a linear function $f(\mathbf{x})$ from a finite data set $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Theorem 1. *Consider the risk minimization problem of the form:*

$$\min_{f \in \mathcal{H}} \underbrace{R_n(\mathbf{y}, \mathbf{f})}_{Empirical\ Risk} + \underbrace{\lambda \Omega(\|f\|_{\mathcal{H}})}_{Regularizer} \tag{40}$$

where $\mathbf{f} = \{f(\mathbf{x}_1), \cdots, f(\mathbf{x}_n)\}, \mathbf{y} = \{y_1, \cdots, y_n\}$, and λ is a scaling parameter. Then 40 always has an optimal solution of the form:

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x}) \tag{41}$$

4 Spectral Kernel

4.1 Fourier Transforms

- Fourier transform $S(\omega)$ of a function $f(x)$,

$$S(\omega) = \int_{-\infty}^{\infty} f(x) \exp(-2\pi i x\omega) dx \tag{42}$$

- Inverse Fourier transform $f(x)$ of a spectral density $S(\omega)$

$$f(x) = \int_{-\infty}^{\infty} S(\omega) \exp(2\pi i x\omega) d\omega \tag{43}$$

- Euler's identity:

$$\exp(ix) = \cos x + i \sin x \tag{44}$$

Hence

$$\exp(\pm 2\pi i x\omega) = \cos(2\pi x\omega) \pm i \sin(2\pi x\omega) \tag{45}$$

4.2 Fourier Duals

Theorem 2. *Bochner's theorem: Any stationary kernel $k : \mathbb{R}^D \to \mathbb{R}$ and its spectral density $S : \mathbb{R}^D \to \mathbb{R}$ are Fourier duals*

$$k(x - x') \equiv k(\tau) = \int_{-\infty}^{\infty} S(\omega) \exp(2\pi i x\omega^\top \tau) d\omega$$
$$S(\omega) = \int_{-\infty}^{\infty} k(\tau) \exp(-2\pi i x\omega^\top \tau) d\tau$$

5 Marginal Likelihood via Laplace Approximation

- Marginal likelihood to do model selection:

$$p(\mathbf{y}) = \int p(\mathbf{y}|\mathbf{f})\, p(\mathbf{f})\, d\mathbf{f} \tag{46}$$

- Let $\psi(\mathbf{f}) = \log h(\mathbf{f}) = \log(p(\mathbf{y}|\mathbf{f})p(\mathbf{f}))$

$$\psi(\mathbf{f}) = \log p(\mathbf{y}|\mathbf{f}) - \frac{N}{2} \log 2\pi - \frac{1}{2} \log |\mathbf{K}| - \frac{1}{2} \mathbf{f}^\top \mathbf{K}^{-1} \mathbf{f} \tag{47}$$

- Second order Taylor approximation around the mode $\hat{\mathbf{f}}$

$$\psi(\mathbf{f}) = \psi(\hat{\mathbf{f}}) - \frac{1}{2} (\mathbf{f} - \hat{\mathbf{f}})^\top \mathbf{A} (\mathbf{f} - \hat{\mathbf{f}}) \tag{48}$$

- Substituting back

$$p(\mathbf{y}) \approx q(\mathbf{y}) = \int \exp(\psi(\hat{\mathbf{f}}) - \frac{1}{2}(\mathbf{f}-\hat{\mathbf{f}})^\top \mathbf{A}(\mathbf{f}-\hat{\mathbf{f}}))d\mathbf{f} \tag{49}$$

$$= \exp(\psi(\hat{\mathbf{f}})) \int \exp(-\frac{1}{2}(\mathbf{f}-\hat{\mathbf{f}})^\top \mathbf{A}(\mathbf{f}-\hat{\mathbf{f}}))d\mathbf{f} \tag{50}$$

$$= \exp(\psi(\hat{\mathbf{f}}))(2\pi)^{N/2}|\mathbf{A}^{-1}|^{1/2} \tag{51}$$

$$= \exp(\log p(\mathbf{y}|\hat{\mathbf{f}}) - \frac{N}{2}\log 2\pi - \frac{1}{2}\log|\mathbf{K}| - \frac{1}{2}\hat{\mathbf{f}}^\top\mathbf{K}^{-1}\hat{\mathbf{f}}) \\ (2\pi)^{N/2}|\mathbf{A}^{-1}|^{1/2} \tag{52}$$

- Taking the log of $q(\mathbf{y})$

$$\log q(\mathbf{y}) = \log p(\mathbf{y}|\hat{\mathbf{f}}) - \frac{N}{2}\log 2\pi - \frac{1}{2}\log|\mathbf{K}| - \frac{1}{2}\hat{\mathbf{f}}^\top\mathbf{K}^{-1}\hat{\mathbf{f}} \\ + \frac{N}{2}\log 2\pi + \frac{1}{2}\log|\mathbf{A}|^{-1} \tag{53}$$

$$= \log p(\mathbf{y}|\hat{f}) - \frac{1}{2}\log|\mathbf{K}| - \frac{1}{2}\hat{\mathbf{f}}^\top\mathbf{K}^{-1}\hat{\mathbf{f}} + \frac{1}{2}|\mathbf{A}^{-1}| \tag{54}$$

- We can now use the fact that $|\mathbf{A}^{-1}| = |\mathbf{A}|^{-1}$

$$\log q(\mathbf{y}) = \log p(\mathbf{y}|\hat{\mathbf{f}}) - \frac{1}{2}\log|\mathbf{K}| - \frac{1}{2}\hat{\mathbf{f}}^\top\mathbf{K}^{-1}\hat{\mathbf{f}} - \frac{1}{2}|\mathbf{A}| \tag{55}$$

- Recall that $\mathbf{A} = \mathbf{K}^{-1} + \mathbf{W}$

$$\log q(\mathbf{y}) = \log p(\mathbf{y}|\hat{\mathbf{f}}) - \frac{1}{2}\log|\mathbf{K}| - \frac{1}{2}\hat{\mathbf{f}}^\top\mathbf{K}^{-1}\hat{\mathbf{f}} - \frac{1}{2}|\mathbf{K}^{-1} + \mathbf{W}| \tag{56}$$

- We optimize $\log q(\mathbf{y})$ using gradient based methods to choose hyperparameters.

6 Multi-output GP

6.1 Intrinsic coregionalization model (ICM): two-outputs

- Consider two output $f_1(\mathbf{x})$ and $f_2(\mathbf{x})$ with $\mathbf{x} \in \mathbb{R}^d$
- Assume the following generative model:
  1. Sample from a GP $u(\mathbf{x}) \sim \mathcal{GP}(0, k(\mathbf{x}, \mathbf{x}'))$ to obtain $u^1(\mathbf{x})$
  2. Obtain $f_1(\mathbf{x})$ and $f_2(\mathbf{x})$ by linearly transforming $u^1(\mathbf{x})$

$$f_1(\mathbf{x}) = a_1^1 u(\mathbf{x})$$
$$f_2(\mathbf{x}) = a_2^1 u(\mathbf{x})$$

6.2 ICM: covariance

- For a fixed value $\mathbf{x}$, we can group $f_1(\mathbf{x})$ and $f_2(\mathbf{x})$ in a vector $\mathbf{f}(\mathbf{x})$

$$\mathbf{f}(\mathbf{x}) = \begin{bmatrix} f_1(\mathbf{x}) \\ f_2(\mathbf{x}) \end{bmatrix}$$

  We refer to this as a vector-valued function.

- The covariance for $\mathbf{f}(\mathbf{x})$ is computed as

$$cov(\mathbf{f}(\mathbf{x}), \mathbf{f}(\mathbf{x})) = \mathbb{E}[\mathbf{f}(\mathbf{x})\mathbf{f}(\mathbf{x}')^\top] - \mathbb{E}[\mathbf{f}(\mathbf{x})]\mathbb{E}[\mathbf{f}(\mathbf{x}')]^\top$$

- We compute the term $\mathbb{E}[\mathbf{f}(\mathbf{x})\mathbf{f}(\mathbf{x}')^\top]$

$$\begin{aligned}
\mathbb{E}\left[\begin{bmatrix} f_1(\mathbf{x}) \\ f_2(\mathbf{x}) \end{bmatrix}\begin{bmatrix} f_1(\mathbf{x}') & f_2(\mathbf{x}') \end{bmatrix}\right] &= \begin{bmatrix} \mathbb{E}[f_1(\mathbf{x})f_1(\mathbf{x}')] & \mathbb{E}[f_1(\mathbf{x})f_2(\mathbf{x}')] \\ \mathbb{E}[f_2(\mathbf{x})f_1(\mathbf{x}')] & \mathbb{E}[f_2(\mathbf{x})f_2(\mathbf{x}')] \end{bmatrix} \\
&= \begin{bmatrix} (a_1^1)^2\mathbb{E}[u_1(\mathbf{x})u^1(\mathbf{x}')] & a_1^1a_2^1\mathbb{E}[u_1(\mathbf{x})u^1(\mathbf{x}')] \\ a_1^1a_2^1\mathbb{E}[u_1(\mathbf{x})u^1(\mathbf{x}')] & (a_2^1)^2\mathbb{E}[u_1(\mathbf{x})u^1(\mathbf{x}')] \end{bmatrix} \\
&= \begin{bmatrix} (a_1^1) & a_1^1a_2^1 \\ a_1^1a_2^1 & (a_2^1)^2 \end{bmatrix}\mathbb{E}[u^1(\mathbf{x})u^1(\mathbf{x}')]
\end{aligned}$$

- The term $\mathbb{E}[\mathbf{f}(\mathbf{x})]$ is computed as

$$\mathbb{E}\left[\begin{bmatrix} f_1(\mathbf{x}) \\ f_2(\mathbf{x}) \end{bmatrix}\right] = \begin{bmatrix} \mathbb{E}[f_1(\mathbf{x})] \\ \mathbb{E}[f_2(\mathbf{x})] \end{bmatrix} = \begin{bmatrix} a_1^1 \\ a_2^1 \end{bmatrix}\mathbb{E}[u^1(\mathbf{x})]$$

- Putting the terms together, the covariance for $\mathbf{f}(\mathbf{x})$ follows

$$\begin{bmatrix} (a_1^1) & a_1^1a_2^1 \\ a_1^1a_2^1 & (a_2^1)^2 \end{bmatrix}\mathbb{E}[u^1(\mathbf{x})u^1(\mathbf{x}')] - \begin{bmatrix} a_1^1 \\ a_2^1 \end{bmatrix}\begin{bmatrix} a_1^1 & a_2^1 \end{bmatrix}\mathbb{E}[u^1(\mathbf{x})]\mathbb{E}[u^1(\mathbf{x}')]$$

- Defining $\mathbf{a} = \begin{bmatrix} a_1^1 & a_2^1 \end{bmatrix}^\top$ and $\mathbf{B} = \mathbf{a}\mathbf{a}^\top$,

$$cov(\mathbf{f}(\mathbf{x}), \mathbf{f}(\mathbf{x}')) = \mathbf{a}\mathbf{a}^\top k(\mathbf{x}, \mathbf{x}') = \mathbf{B}^\top k(\mathbf{x}, \mathbf{x}')$$

6.3 ICM: Observed data

- Given $\mathcal{D}_1 = \{(\mathbf{x}_i, f_1(\mathbf{x}_i))|i = 1, \cdots, N\}$ and $\mathcal{D}_2 = \{(\mathbf{x}_i, f_2(\mathbf{x}_i))|i = 1, \cdots, N\}$, then

$$\begin{bmatrix} \mathbf{f}_1 \\ \mathbf{f}_2 \end{bmatrix} = \begin{bmatrix} f_1(\mathbf{x}_1) \\ \vdots \\ f_1(\mathbf{x}_N) \\ f_2(\mathbf{x}_1) \\ \vdots \\ f_2(\mathbf{x}_N) \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} b_{11}\mathbf{K} & b_{12}\mathbf{K} \\ b_{21}\mathbf{K} & b_{22}\mathbf{K} \end{bmatrix}\right) = \mathcal{N}(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \mathbf{B} \otimes \mathbf{K})$$

- The inversion rule: $(A \otimes B)^{-1} = A^{-1} \otimes B^{-1}$

7 Computational Complexity of GP Regression

- Data set with N observations, computing posterior for 1 test point:

$$\mu_* = \mathbf{K}_{f_*f}(\mathbf{K}_{ff} + \sigma^2\mathbf{I})^{-1}\mathbf{y}$$
$$\sigma_*^2 = \mathbf{K}_{f_*f_*} - \mathbf{K}_{f_*f}(\mathbf{K}_{ff} + \sigma^2 I)^{-1}\mathbf{K}_{f_*f}^\top$$

- Matrix-vector multiplication (mvm): for $\mathbf{A} \in \mathbb{R}^{N\times M}$ and $\mathbf{b} \in \mathbb{R}^M$, computing $\mathbf{Ab}$ costs $\mathcal{O}(NM)$
- Matrix inverse: for $\mathbf{C} \in \mathbb{R}^{N\times N}$, computing $\mathbf{C}^{-1}$ costs $\mathcal{O}(N^3)$
- $(\mathbf{K}_{ff} + \sigma^2\mathbf{I})^{-1}\mathbf{y}$ scales as $\mathcal{O}(N^3)$.

8 Approximately solving linear system

8.1 Matrix inverse as quadratic optimization

- Rewrite matrix inverse

$$\mathbf{v} = \hat{\mathbf{K}}^{-1}\mathbf{y}, \qquad \hat{\mathbf{K}} = \mathbf{K} + \sigma^2\mathbf{I}$$

  as a linear system:

$$\hat{\mathbf{K}}\mathbf{v} - \mathbf{y} = 0$$

- Solve as a quadratic optimization problem:

$$\mathbf{v}^* = \arg\min_{\mathbf{v}} \mathbf{v}^\top\hat{\mathbf{K}}\mathbf{v} - \mathbf{v}^\top\mathbf{y}$$

8.2 Conjugate gradient

- Using conjugate gradient to solve the quadratic optimization
  1. Iterative method
  2. Each step is $\mathcal{O}(N^2)$
  3. Recovers exact solution after N steps $\rightarrow \mathcal{O}(N^3)$
  4. Approximate solution in much fewer steps: less steps.

8.3 Convergence and preconditioning

- Condition number: ratio of largest to smallest eigenvalue $\lambda_{\min}(\hat{\mathbf{K}})/\lambda_{\max}(\hat{\mathbf{K}})$.
- High condition numbers: numerically unstable, slow convergence.
- Improve by preconditioning: Instead of $\hat{\mathbf{K}}\mathbf{v} - \mathbf{y} = 0$, solve

$$\mathbf{P}^{-1}\hat{\mathbf{K}}\mathbf{v} - \mathbf{P}^{-1}\mathbf{y} = 0$$

9 Low-rank approximation

- Recall GP marginal log-likelihood:

$$\log p(\mathbf{y}|\mathbf{X}) = \log \mathcal{N}(\mathbf{y}|0, \mathbf{K} + \sigma^2\mathbf{I})$$

  Assume $\mathbf{K}$ to be low rank.

9.1 Approximation by subset

- Let's randomly pick a subset from training data: $\mathbf{Z} \in \mathbb{R}^{M\times Q}$
- Approximate the covariance matrix $\mathbf{K}$ by $\hat{K}$

$$\hat{\mathbf{K}} = \mathbf{K}_z\mathbf{K}_{zz}^{-1}\mathbf{K}_z^\top \in \mathbb{R}^{N\times N}$$
$$\mathbf{K}_z = \mathbf{K}(\mathbf{X}, \mathbf{Z}) \in \mathbb{R}^{N\times M}$$
$$\mathbf{K}_{zz} = \mathbf{K}(\mathbf{Z}, \mathbf{Z}) \in \mathbb{R}^{M\times M}$$

- The log-likelihood is approximated by

$$\log p(\mathbf{y}|\mathbf{X}) = \log \mathcal{N}(\mathbf{y}|0, \mathbf{K}_z\mathbf{K}_{zz}^{-1}\mathbf{K}_z^\top + \sigma^2\mathbf{I})$$

- Furthermore, apply Woodbury matrix identity:

$$(UCV + A)^{-1} = A^{-1} - A^{-1}U(C^{-1} + VA^{-1}U)^{-1}VA^{-1}$$
$$(\mathbf{K}_z\mathbf{K}_{zz}^{-1}\mathbf{K}_z^\top + \sigma^2\mathbf{I})^{-1} = \sigma^{-2}\mathbf{I} - \sigma^{-4}\mathbf{K}_z(\mathbf{K}_{zz} + \sigma^{-2}\mathbf{K}_z^\top\mathbf{K}_z)^{-1}\mathbf{K}_z^\top$$

- The complexity reduces to $\mathcal{O}(NM^2)$.

10 Variational Inference for Sparse GP

10.1 Inducing point methods: the joint model

- Goal: choose a set of inducing points s.t. it contains the same information as a full data set.
- The augmented model

$$p(\mathbf{y}, \mathbf{f}, \mathbf{u}) = p(\mathbf{y}|\mathbf{f})p(\mathbf{f}, \mathbf{u})$$

- Recover the original model by marginalizing over $\mathbf{u}$:

$$p(\mathbf{y}, \mathbf{f}) = \int p(\mathbf{y}|\mathbf{f})p(\mathbf{f}, \mathbf{u})d\mathbf{u} = p(\mathbf{y}|\mathbf{u}) \int p(\mathbf{f}, \mathbf{u})d\mathbf{u} = p(\mathbf{y}|\mathbf{f})p(\mathbf{f}) = p(\mathbf{y}|\mathbf{f})p(\mathbf{f}|\mathbf{u})p(\mathbf{u})$$

- Using Gaussian conditional densities:

$$p(\mathbf{y}|\mathbf{f}) = \mathcal{N}(\mathbf{y}|\mathbf{f}, \sigma^2\mathbf{I})$$
$$p(\mathbf{f}|\mathbf{u}) = \mathcal{N}(\mathbf{f}|\mathbf{K}_{nm}\mathbf{K}_{mm}^{-1}\mathbf{u}, \hat{\mathbf{K}}), \quad \hat{K} = \mathbf{K}_{nn} - \mathbf{K}_{nm}\mathbf{K}_{mm}^{-1}\mathbf{K}_{mn}$$
$$p(\mathbf{u}) = \mathcal{N}(\mathbf{u}|0, \mathbf{K}_{mm})$$

- Covariance of inducing points: $[\mathbf{K}_{mm}]_{ij} = k(z_i, z_j)$
- Cross-covariance between inducing points and training: $[\mathbf{K}_{mn}]_{ij} = k(z_i, x_j)$
- Covariance of training points: $[\mathbf{K}_{nn}]_{ij} = k(x_i, x_j)$

10.2 Variational Sparse GP

- Variational lower bound of a marginal likelihood:

$$\begin{aligned}\log p(\mathbf{y}|\mathbf{X}) &= \log \int_{\mathbf{f},\mathbf{u}} p(\mathbf{y}|\mathbf{f})p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})p(\mathbf{u}|\mathbf{Z}) \\ &\geq \int_{\mathbf{f},\mathbf{u}} q(\mathbf{f},\mathbf{u}) \log \frac{p(\mathbf{y}|\mathbf{f})p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})p(\mathbf{u}|\mathbf{Z})}{q(\mathbf{f},\mathbf{u})} \equiv \mathcal{L}\end{aligned}$$

- Defining the variational posterior as follows:

$$\begin{aligned}q(\mathbf{f},\mathbf{u}) &= p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})q(\mathbf{u}) \\ q(\mathbf{u}) &= \mathcal{N}(\mathbf{u}|\mu,\Sigma)\end{aligned}$$

- Therefore, we have

$$\begin{aligned}\mathcal{L} &= \int_{\mathbf{f},\mathbf{u}} p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})q(\mathbf{u}) \log \frac{p(\mathbf{y}|\mathbf{f})\cancel{p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})}p(\mathbf{u}|\mathbf{Z})}{\cancel{p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})}q(\mathbf{u})} \\ &= \langle \log p(\mathbf{y}|\mathbf{f}) \rangle_{p(\mathbf{f}|\mathbf{u},\mathbf{X},\mathbf{Z})q(\mathbf{u})} - \mathrm{KL}[q(\mathbf{u})|p(\mathbf{u}|\mathbf{Z})]\end{aligned}$$